



## Lab: Code Safety: Types

Software, hardware, and systems should do everything that they are intended to do and nothing more. The “nothing more” part of that statement is what is meant by safety. If, say code, is allowed to do more than what is intended an attacker just might be able to exploit unintended execution of given code and perform some malicious task or compromise privacy, for example. Numerous programming languages require great care to develop a safe product. The C/C++ family is most notorious in allowing a developer to create exploitable code but other popular languages do so as well. Three classes of safety are of particular importance because they have been most responsible for unintended execution. These are: memory safety, thread safety, and type safety. This lab is concerned with type safety. By type safety we mean:

Type Safety:

1. No undesirable program behavior that is caused by a discrepancy between differing data types for constants, variables, and functions. Types could be manufactured (structs, classes) or primitive.
2. Integer overflow, underflow, casting should not cause undesirable program behavior.

Consider the following C++ code, provided in file `linear.c`, for illustration:

```
#include <stdlib.h>
#include <stdio.h>
#include <stdint.h>

// Input should be increasing set of distinct numbers
uint8_t inputOK (uint16_t elem[]) {
    int i;
    for (i=0 ; i < 232 ; i++) if (elem[i] >= elem[i+1]) return 0;
    return 1;
}

uint8_t lin_search (uint16_t key, uint16_t xs[]) {
    uint8_t i;
    if (inputOK(xs) == 0) return -1;
    for (i=0 ; i < 233 ; i++) if (key == xs[i]) return i;
    return 255;
}
```

The problem solved by function `lin_search` is: given a list (or sequence) `xs` of 233 positive distinct 16 bit integers in increasing order and a positive 16 bit integer `key`, find the position of `key` in `xs`. If `key` is not a member of `xs` then -1 is returned. For example, given `key` is 4 and `xs` is 1,2,3,4,5,6,7 then the position of `key` in `xs` is 3. Function `inputOK` is used to ensure integers in `xs` are distinct and increasing in order. If not, -1 is returned.

In file `linear_search.cry` Cryptol function `linSearchIdx` returns the position of a given `key` in given sequence `xs` or -1. Function `linSearchIdx` is simple: a sequence `s`, initialized with

a single -1, is extended by comparing key with consecutive elements  $x$  of  $xs$  (via  $x \leftarrow xs$ ). A counter  $c$  keeps track of the position in  $xs$  of element  $x$ . If key is not equal to  $x$ , then the next element in sequence  $s$  is the same as the previous element. If key has not yet been found in  $xs$ , that will be -1, but if key is equal to  $x$  the next element in  $s$  is the value of  $c$ . The purpose of  $h$  is to ensure that  $s$  contains the value of  $c$  when key equals  $x$  for all of  $s$  after equality is established (otherwise  $s$  is all -1s). The output of `linSearchIdx` is the value of the last element of  $s$  if `inputOK`, which has the same functionality as its namesake in the C code, returns True on  $xs$  or -1 if not. The signature of `linSearchIdx` is monomorphic: the key is 16 bits, each of 233 integers in  $xs$  are 16 bits and the position that is output is 8 bits. These numbers match those for the C code. Here is the Cryptol specification for `inputOK` and `linSearchIdx`:

```
// Returns True iff elmts is a sequence whose numbers are always increasing
// in order (all are distinct)
inputOK : [233][16] -> [8]
inputOK elmts = z ! 0
  where
    z = [1]#[ if x >= y /\ q == 1 then 0 else q
              | q <- z | x <- elmts | y <- (tail elmts) ]

// Returns the first location of a given key in the given list
// or -1 if the key is not in the list
linSearchIdx : [16] -> [233][16] -> [8]
linSearchIdx key xs = if (inputOK xs) == 1 then s!0 else -1
  where
    s = [255]#[ if (key == x /\ h == 255) then c else h
                | x <- xs | h <- s | c <- [0...]]
```

**Note:** SAW doesn't yet support translating Cryptol's bit type(s) into crucible-llvm's type system so the output type of `inputOK` is `[8]` instead of `Bit`.

A SAW file `a.saw` is created to check the equivalence of `lin_search` and `linSearchIdx`. This requires LLVM bitcode for functions in file `linear.c` which is obtained from this:

```
clang -g -O0 -emit-llvm -c linear.c -o linear.bc
```

File `a.saw` follows patterns established in the section on memory safety:

```
import "linear_search.cry";

let safe_setup = do {
  xs <- llvm_fresh_var "array" (llvm_array 233 (llvm_int 16));
  pxs <- llvm_alloc (llvm_array 233 (llvm_int 16));
  llvm_points_to pxs (llvm_term xs);
  key <- llvm_fresh_var "key" (llvm_int 16);
  llvm_execute_func [ llvm_term key, pxs ];
  llvm_return (llvm_term [{ linSearchIdx key xs }]);
};

let main : TopLevel () = do {
  m <- llvm_load_module "linear.bc";
  saf_proof <- llvm_verify m "lin_search" [] true safe_setup abc;
  print "Done!";
};
```

Run `saw a.saw` to get this:

```
[13:35:52.130] Verifying lin_search ...
[13:35:52.140] Simulating lin_search ...
[13:35:59.688] Checking proof obligations lin_search ...
[13:36:00.684] Proof succeeded! lin_search
[13:36:00.684] Done!
```

The problem of finding the position of a key in an increasing sequence of distinct integers is normally solved using binary search. A common way this is presented is shown in file `bsearch-unsafe.c`:

```
#include <stdlib.h>
#include <stdint.h>
#include <stdio.h>

// Input should be increasing set of distinct numbers
uint8_t inputOK (uint16_t elem[]) {
    int i;
    for (i=0 ; i < 232 ; i++) if (elem[i] >= elem[i+1]) return 0;
    return 1;
}

// Standard binary search except index can overflow for high enough f and l
uint8_t binsearch (uint8_t f, uint8_t l, uint16_t x, uint16_t elem[]) {
    uint8_t index = 0;
    while (f <= l) {
        index = (f+l);
        index = index/2;
        if (x == elem[index]) return index;
        if (x < elem[index]) l = index-1;
        else f = index+1;
    }
    return 255;
}

// Call binary search, assume list size is 233
uint8_t b_search (uint16_t key, uint16_t xs[]) {
    if (inputOK(xs) == 1) return binsearch(0, 232, key, xs);
    else return 255;
}
```

The reason for `index = (f+l)`; followed by `index = index/2`; instead of `index = (f+l)/2`; is that some compilers will generate code to prevent the overflow and, as noted earlier, this overflow is a problem and is not usually solved in languages such as C. The above code forces the overflow.

For example, add the following main to the file:

```
int main (int argc, char **argv) {
    uint16_t elems[] = {
        23, 26, 33, 40, 49, 50, 53, 60, 62, 67, 71, 77, 80, 85, 94,
        103, 109, 116, 117, 123, 124, 125, 126, 134, 135, 139, 140, 141, 150, 152,
        157, 165, 171, 173, 179, 184, 188, 197, 199, 206, 211, 218, 222, 230, 232,
        235, 243, 252, 253, 261, 265, 266, 267, 273, 281, 284, 285, 286, 288, 289,
        290, 299, 308, 316, 318, 323, 326, 333, 338, 344, 349, 350, 353, 362, 371,
        378, 380, 387, 393, 396, 402, 403, 407, 415, 423, 426, 428, 436, 439, 445,
        454, 459, 465, 473, 477, 486, 489, 497, 503, 512, 516, 518, 527, 536, 537,
        545, 551, 555, 561, 563, 570, 573, 577, 579, 580, 582, 587, 591, 592, 601,
```

```

        602, 603, 607, 614, 623, 632, 638, 641, 648, 651, 653, 655, 660, 662, 663,
        670, 671, 678, 679, 686, 687, 695, 696, 700, 701, 703, 710, 717, 723, 732,
        738, 744, 745, 746, 750, 751, 752, 753, 757, 765, 770, 776, 777, 778, 787,
        789, 797, 798, 807, 808, 815, 816, 817, 824, 829, 830, 831, 834, 843, 851,
        854, 859, 864, 868, 874, 875, 879, 886, 889, 898, 903, 911, 918, 924, 925,
        931, 939, 948, 956, 964, 973, 980, 989, 998, 1004, 1008, 1009, 1017, 1025,
        1031, 1033, 1039, 1041, 1048, 1051, 1054, 1056, 1057, 1062, 1064, 1070,
        1075, 1084, 1086, 1091, 1098, 1099, 1103, 1108, 1117, 1121, 1124, 1130 };
    printf("%d\n", b_search(atoi(argv[1]), elems));
}

```

Then run

```

[prompt]$ bsearch-unsafe 165
31
[prompt]$ bsearch-unsafe 587
116
[prompt]$ bsearch-unsafe 591
117
[prompt]$ bsearch-unsafe 1117
infinite loop

```

The last error is due to integer overflow in  $(f+l)$  when computing the midpoint. The following change to `binsearch`, in file `bsearch-safe.c`, uses a midpoint calculation that does not overflow:

```

uint8_t binsearch (uint8_t f, uint8_t l, uint16_t x, uint16_t elem[]) {
    uint8_t index = 0;
    while (f <= l) {
        index = (f+(l-f)/2);
        if (x > elem[index])
            f = index+1;
        else if (x < elem[index])
            l = index-1;
        else return index;
    }
    return 255;
}

```

Compile `bsearch-safe.c` and run as follows:

```

[prompt]$ bsearch-unsafe 165
31
[prompt]$ bsearch-unsafe 587
116
[prompt]$ bsearch-unsafe 591
117
[prompt]$ bsearch-unsafe 1117
229

```

Next, it is instructive to use SAW to show `b_search` in `bsearch-unsafe.c` is not safe and `b_search` in `bsearch-safe.c` is equivalent to the search specification which is safe. Run Cryptol and load `linear_search.cry`. Then run `:safe linSearchIdx` and `:safe inputOK` to show that `linSearchIdx` is safe. Create LLVM bitcode for functions in `bsearch-unsafe.c` and `bsearch-safe.c` with this:

```

clang -g -O0 -emit-llvm -c bsearch-unsafe.c -o bsearch_unsafe.bc
clang -g -O0 -emit-llvm -c bsearch-safe.c -o bsearch_safe.bc

```

The SAW file for `bsearch-safe.c` is `b.saw` and looks like this:

```
import "linear_search.cry";

let safe_setup = do {
  xs <- llvm_fresh_var "array" (llvm_array 233 (llvm_int 16));
  pxs <- llvm_alloc (llvm_array 233 (llvm_int 16));
  llvm_points_to pxs (llvm_term xs);
  key <- llvm_fresh_var "key" (llvm_int 16);
  llvm_execute_func [ llvm_term key, pxs ];
  llvm_return (llvm_term {{ linSearchIdx key xs }});
};

let inputOK_setup = do {
  xs <- llvm_fresh_var "array" (llvm_array 233 (llvm_int 16));
  pxs <- llvm_alloc (llvm_array 233 (llvm_int 16));
  llvm_points_to pxs (llvm_term xs);
  llvm_execute_func [ pxs ];
  llvm_return (llvm_term {{ inputOK xs }});
};

let main : TopLevel () = do {
  m <- llvm_load_module "bsearch_safe.bc";
  inp_OK <- llvm_verify m "inputOK" [] true inputOK_setup abc;
  saf_proof <- llvm_verify m "b_search" [inp_OK] true safe_setup abc;
  print "Done!";
};
```

We will skip doing this for `bsearch-unsafe.c` as this will not add anything meaningful to the discussion.

Run `saw b.saw` to get an error that is listed above a ‘counterexample’:

```
[14:13:28.707] Subgoal failed: b_search safety assertion:
bsearch-safe.c:17:16: error: in binsearch
Error during memory load
```

The problem occurs because when SAW replaces the call to `inputOK` in `b_search` with the Cryptol specification, SAW must conservatively consider the possibility that `inputOK` invalidated the `xs` pointer. To fix this, `llvm_alloc_readonly` is used instead of `llvm_alloc` so SAW knows `inputOK` does not modify `xs`. Observe `a.saw` succeeded but uses `llvm_alloc` when it should have used `llvm_alloc_readonly`, which is confusing. Better safe than sorry and use the `readonly` allocation whenever functions do not change input structures.

Running SAW with the above change to `llvm` allocation takes a lot of time and likely does not finish. Since this is intended as an illustration of how to use SAW, consider modifying `b_search`, `inputOK`, `linSearchIdx`, and `b.saw` for lists of 10 elements instead of 233. Here is a guide for making those changes, make them in files `bsearch-safe-1.c`, `bb.saw`, and `linear_search_1.cry`:

```
C (bsearch-safe-1.c):
inputOK: use for (i=0 ; i < 9 ; i++)
b_search: use return binsearch(0, 9, key, xs);
Cryptol (linear_search_1.cry):
inputOK : [10][16] -> [8]
```

```
linSearchIdx : [16] -> [10][16] -> [8]
```

bb.saw:

```
import "linear_search_1.cry";
safe_setup:
  xs <- llvm_fresh_var "array" (llvm_array 10 (llvm_int 16));
  pxs <- llvm_alloc_readonly (llvm_array 10 (llvm_int 16));
inputOK_setup:
  xs <- llvm_fresh_var "array" (llvm_array 10 (llvm_int 16));
  pxs <- llvm_alloc_readonly (llvm_array 10 (llvm_int 16));
main:
  m <- llvm_load_module "bsearch_safe_1.bc";
```

Run

```
clang -g -O0 -emit-llvm -c bsearch-safe-1.c -o bsearch_safe_1.bc
```

and then saw bb.saw to get this:

```
[15:25:45.167] Subgoal failed: b_search safety assertion:
bsearch-safe-1.c:20:11: error: in binsearch
Undefined behavior encountered
Details:
  Addition of an offset to a pointer resulted in a pointer to an address
  outside of the allocation
```

So, a problem with the C code has been identified: the binsearch implementation contains an out-of-bounds memory access when key is less than the smallest element in xs. This is because index eventually underflows and so elem[index] is out-of-bounds. One fix is this for binsearch:

```
uint8_t binsearch (uint8_t f, uint8_t l, uint16_t x, uint16_t elem[]) {
  uint8_t index = 0;
  while (f <= l) {
    index = (f+(l-f)/2);
    if (index >= 0 && index <= 9 && x == elem[index]) return index;
    if (index >= 0 && index <= 9 && x < elem[index]) l = index-1;
    else if (index >= 0 && index <= 9) f = index+1;
    else return 255;
  }
  return 255;
}
```

Let file bsearch-safe-2.c be bsearch-safe-1.c with this fix. Then

```
clang -g -O0 -emit-llvm -c bsearch-safe-2.c -o bsearch_safe_2.bc
```

plus copy bb.saw to bbb.saw and change bsearch\_safe\_1.bc to bsearch\_safe\_2.bc.

Run saw bbb.saw to get:

```
[16:47:57.463] Verifying inputOK ...
[16:47:57.465] Simulating inputOK ...
[16:47:57.489] Checking proof obligations inputOK ...
[16:47:57.540] Proof succeeded! inputOK
[16:47:57.575] Verifying b_search ...
[16:47:57.576] Simulating b_search ...
[16:47:57.578] Registering overrides for `inputOK`
[16:47:57.578]   variant `Symbol "inputOK"`
[16:47:57.578] Matching 1 overrides of inputOK ...
[16:47:57.578] Branching on 1 override variants of inputOK ...
[16:47:57.578] Applied override! inputOK
```

```
[16:47:57.630] Checking proof obligations b_search ...  
[16:47:58.036] Proof succeeded! b_search  
[16:47:58.036] Done!
```